

ML Learning Assignment Project Report



Dhruv Maheshwari - PES2UG23CS173

Sem 5 – C

Date: 12th Oct 2025

Divyansh Sharma - PES2UG23CS181

Project ID: 14

Team ID: 14

Project Title: Intraday Stock Price Forecasting with Bidirectional LSTM

Problem Statement

The main problem addressed is intraday stock price forecasting for a given ticker (GOOG in the example notebook) using high-frequency (5-minute interval) data. The goal is to predict short-term price movements to support potential trading decisions, which is challenging due to the inherent volatility and non-linear nature of financial time series.

Objective

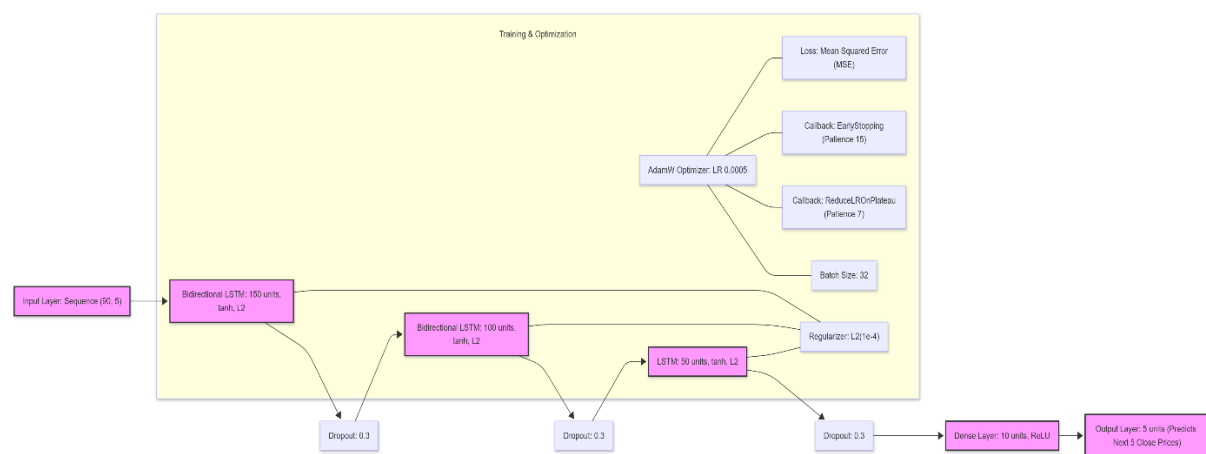
The project aims to develop an optimized and robust Bidirectional Long Short-Term Memory (Bi-LSTM) neural network model capable of multi-step prediction to forecast the Close price for the next 5 five-minute intervals based on the stock's recent historical data.

Dataset Details

Detail	Value
Source	yfinance library (Yahoo)
Size	1673 samples (1 month - 5min intervals)
Key Features	Open, High, Low, Close, Volume
Target Variable	Future Close Price for next 5 intervals

Architecture Diagram

The architecture is a Deep Bidirectional LSTM network.



Layer Structure

- Input Layer: Expects a sequence of 90 time steps and 5 features.
- Bidirectional LSTM (150 units): Allows the network to process the sequence both forward and backward for better context capture. Uses tanh activation and L2 regularization.
- Dropout (0.3): Applied for regularization to prevent overfitting.
- Bidirectional LSTM (100 units): Second Bi-LSTM layer, also with tanh activation and L2 regularization.
- LSTM (50 units): Standard LSTM layer, returning only the final sequence output.
- Dense Layer (10 units): Intermediate dense layer with relu activation.
- Output Layer (5 units): A Dense layer predicting the 5-step future Close prices.

Methodology

The workflow employed a robust, step-by-step process:

- Data Acquisition: GOOG (Alphabet Inc.) stock price data was fetched for 1 month at 5-minute intervals using the yfinance library.
- Feature Selection: The 'Open, High, Low, Close, Volume' features were selected.
- Scaling: The data was scaled using RobustScaler to minimize the impact of potential outliers, common in financial data.
- Sequence Creation & Splitting: Data was transformed into sequences with a length of 90-time steps to predict a horizon of 5 future steps (multi-step prediction). The dataset was split into 80% training and 20% testing sets.
- Model Configuration: A deep Bi-LSTM model was compiled using the AdamW optimizer (with a learning rate of 0.0005) and Mean Squared Error (MSE) as the loss function. L2 regularization was applied to the LSTM layers.
- Training: The model was trained for up to 100 epochs with a small batch size of 32. EarlyStopping and ReduceLROnPlateau callbacks were used to optimize convergence and prevent overfitting.

Results & Evaluation

The model was evaluated using the 1-Step Ahead prediction (the first value in the 5-step output) against the corresponding actual value.

Metric	Value
Mean Squared Error (MSE)	0.3829
Root Mean Squared Error (RMSE)	0.6188
Mean Absolute Error (MAE)	0.4336
R ² Score	0.9741

The R² Score of 0.9741 suggests the model's predictions closely track the actual price movements in the test set. An RMSE of 0.6188 indicates that the average prediction error, in unscaled USD, is approximately 62 cents for the 1-step ahead forecast.

Conclusion

The deep Bidirectional LSTM model successfully learned the sequential patterns in the intraday stock data. The high R^2 score demonstrates strong predictive power for short-term stock price movements. The use of RobustScaler and Bidirectional layers appears to have contributed to the model's performance by making it less sensitive to outliers and better at capturing contextual dependencies in the time series. A key learning point was that the multi-step block prediction approach efficiently generates longer-term forecasts (e.g., the next 30 intervals) by reducing the number of intensive model calls.
